

Run Forensic Audit

Task 1e67725f — “Find 10 qualified AI startups in the US currently hiring”

Every node the run executed, what each decided, the exact actor inputs, and the point at which it stopped — reconstructed from edge-function logs, `tool_calls` and the task row.

Field	Value
Task	1e67725f-bd24-4d62-b76b-012373341ed6
Plan	3bb9d5e3-2fdd-4da5-8f6a-9eeba9826bc5
Started	2026-08-20 08:56:30Z
Last activity	2026-08-20 08:58:01Z
Worker killed	2026-08-20 08:58:56Z — shutdown
Lifetime	2m 26s
Build	run-agent v22 · commit 4c8bdac4 (deployed 08:54:50Z)
Status now	running — and it will never change
Checkpoint	version 0 — none written

Decision owners: **GPT** decides, **CODE** validates and executes, **PROVIDER** returns facts.

1. Verdict

The run was killed mid-qualification, and left no way back.

It reached 3 enriched, eligible companies at 08:58:01 and began GPT qualification. Fifty-five seconds later the edge worker was terminated. Because it died before any terminal or checkpoint write, the task still reads **running, checkpoint_version 0**, and no continuation was ever decided or dispatched. The execution card spins because the row genuinely says the run is still going.

This is NOT the continuation bug fixed in 4c8bdac4.

That defect was a successor *erasing* a checkpoint that existed. Here **no checkpoint was ever written** — the run never got far enough. The fix deployed at 08:54:50 is untested by this run; it was never reached.

Why the existing deadline guard did not save it

`shouldCheckpoint / CHECKPOINT_RESERVE_MS` is consulted inside `callProvider` — before **provider** calls. GPT qualification calls are not provider calls, so nothing checked the clock during the one phase that overran. The reserve protects Apify spend, not wall-clock generally — and qualification is precisely where an unbounded number of model calls happens.

2. Node-by-node timeline

Time	Node	Owner	What happened
08:56:26	worker boot	—	orchestrate booted (91ms)
08:56:27	plan created	CODE	plan 3bb9d5e3 — “Source → rank”
08:56:29	worker boot	—	run-agent booted (95ms)
08:56:30	mission load	GPT	has_mission true · entry startup_company_discovery · output qualified_companies
08:56:30	playbook	CODE	hiring:supported · runnable · nothing blocked
08:56:30	triage binding	CODE	enabled · 4 batches × 25 · pool 100 · cap 6
08:56:30	eval binding	CODE	enabled · 10 calls allowed · shortlist 10 · cap 20
08:56:30	paid preflight	CODE	ok: true
08:56:30	qualification ctx	CODE	required_titles [ai startup, currently hiring]
08:56:35	execution plan	GPT	gpt-4.1 · reasoning · 4,535ms
08:56:35	discovery strategy	CODE	model_repaired · violations [filter_dropped]
08:56:35	memo23	PROVIDE R	→ 100 rows
08:56:50	prequalification	CODE	unique 100 · eligible 0 · technical_only 65
08:57:17– 29	triage ×4	GPT	4 batches × 25 companies
08:57:34	plan amendment	GPT	gpt-4.1 · reasoning · 4,317ms
08:57:34– 42	identity ×8	PROVIDE R	8 company-search calls
08:57:50	identity done	CODE	targets 10 · resolved 3 · unresolved 5 · deferred 2 · truncated
08:57:51	enrichment	PROVIDE R	3 URLs → 3 rows
08:58:01	eligible pool	CODE	discovered 3 · eligible 3
08:58:01	← LAST LOG	GPT	qualification of 3 companies begins
08:58:56	shutdown	infra	worker terminated — no further output
08:58:59	shutdown	infra	second worker

3. What GPT decided, and on what

3.1 Mission (compiled upstream in pilot-chat)

```
entry_capability      startup_company_discovery
requested_output      qualified_companies
verticals             [ai, startup]          stages [startup]
required_titles       ["ai startup", "currently hiring"]
mission_version       lead-mission-v1      mission_type company_research
```

3.2 Execution plan — gpt-4.1, reasoning tier, 4,535ms

```
1 startup_company_discovery      apify_yc_companies_memo23
2 company_identity_resolution    apify_linkedin_company_search
3 company_enrichment             apify_linkedin_company_details
4 company_brain_qualification    (no actor)
5 persistence                    (no actor)

hiring_verification NOT planned -- memo23 carries openJobs
```

3.3 Amendment after discovery — gpt-4.1, 4,317ms

The execute→inspect→replan loop working live. The engine handed GPT a factual pool summary and it re-planned against what actually came back:

```
execution_plan_amended_after_discovery
  observed: [ "10 of 100 rows carry no open role,
              and the mission requires hiring evidence" ]
```

4. Every actor call, with its input

4.1 Discovery — memo23/y-combinator-scraper → 100 rows

```
{
  "mode": "companies",
  "regions": ["United States of America"],
  "isHiring": true,
  "maxItems": 100,
  "industries": ["B2B", "Engineering, Product and Design", "All industries"],
  "enrichEmails": false,
  "scrapeOpenJobs": true,
  "scrapeFounderDetails": false
}
```

Correctly typed arrays, valid enum values, 100 rows on the first attempt — the actor-contract work holding. One incoherence: "All industries" sent alongside two specific buckets.

4.2 Identity resolution — 8 × linkedin-company-search

Each call: {maxItems: 5, scraperMode: "full", searchQuery: "<name>"}

searchQuery	Rows	Outcome
Polymath	5	resolved candidates returned
Deepgram	0	name search found nothing
Rhizome AI	2	resolved → enriched
Robby	5	not resolved to an actionable identity

searchQuery	Rows	Outcome
Parameter	5	not resolved to an actionable identity
Salus	5	not resolved to an actionable identity
RamAln	5	resolved → enriched
Piris Labs	1	resolved → enriched
10 targeted	8 ran	2 deferred on wall clock

4.3 Enrichment — 1 × linkedin-company → 3 rows

```
{ "companies": [  
  "https://www.linkedin.com/company/rhizomeai",  
  "https://www.linkedin.com/company/ramain",  
  "https://www.linkedin.com/company/pirislabs"  
] }
```

5. The funnel

```
100 discovered          memo23, one call
|
0 prequalified          free pass excluded ALL of them
|                       65 as "technical_only"    <-- see 6.1
10 shortlisted          GPT triage overrode the pre-pass
|
8 identity attempted    2 deferred on wall clock
3 resolved              5 unresolved (name search found nothing usable)
|
3 enriched              1 batched call, 3/3 success
3 eligible              ready for qualification
|
? QUALIFIED             <-- worker killed here. Never answered.
```

6. Defects this run exposed

6.1 The free pre-pass disqualifies the mission's own target

```
prequalification_complete { unique: 100, eligible: 0, technical_only: 65 }
```

Every one of the 100 companies was excluded, 65 of them as **technical_only** — meaning they are hiring engineers. For a request asking for *AI startups currently hiring*, disqualifying a company **because** it is hiring engineers is backwards. GPT triage overrode it and shortlisted 10, so it did not sink the run, but the commercial vocabulary is arguing with the mission rather than serving it.

6.2 The observability ledger fails silently, eight times

```
[execution-ledger] insert error [object Object]      x8
```

Every ledger insert failed, and the error was stringified to `[object Object]` — so the reason is unrecoverable from the logs. This is why `lead_execution_calls` holds zero rows.

6.3 A contradictory industries filter was accepted

"All industries" alongside "B2B" and "Engineering, Product and Design". It returned 100 rows so it did no harm here, but the values are mutually exclusive and the card does not say so.

6.4 Cosmetic: the binding logs a model that did not run

Triage and evaluation bindings log model: "openai/gpt-5.6-luna", which is the legacy env diagnostic. The calls actually ran on gpt-4.1, as the `gpt_call_ok` lines show. Confusing during an incident; not a defect in behaviour.

7. What worked

Component	Evidence
Execution planner	Ran, was obeyed, and correctly skipped paid hiring verification because memo23 carries openJobs.
Amendment loop	Fired after discovery, observed the real gap, re-planned. First live proof of execute→inspect→replan.
Model routing	execution_plan and amendment both gpt-4.1/reasoning with stated reasons; triage on the fast tier.
Actor contracts	memo23 input correctly typed with valid enums — 100 rows first attempt, no recovery round needed.
Validator	Caught an unsupported filter (filter_dropped) and repaired the plan rather than failing.

8. The fix this needs

A **wall-clock reserve around qualification**, mirroring the guard that already protects provider calls. Before each evaluation call, check the deadline; when the reserve is reached, stop evaluating, write the checkpoint with whatever has been decided, and let the continuation carry the rest.

```
for (const company of eligible) {  
  if (deadlineReserveReached()) {  
    writeCheckpoint(decidedSoFar); // <-- the safe stopping point  
    return continuationRequired;  //      that does not exist today  
  }  
  await evaluate(company);  
}
```

Today there is no safe stopping point between “enrichment done” and “everything qualified”. For a 10-lead request that is a window the run cannot reliably fit inside one invocation — so this will recur whenever qualification has more than a couple of companies to judge.

Nothing was changed in response to this run. The task remains at running and should be treated as dead; a fresh request is the only way forward until the reserve exists.